

STAT 5244 – Unsupervised Learning

Homework 1

Name: Chuyang Su UNI: cs4570

1 Dimension Reduction on Digits Data.

1.1 Apply linear dimension reduction techniques.

In this experiment, I applied three linear dimension reduction methods and compared their performance on the `scikit-learn Digits` dataset ($n = 1797$, $p = 64$).

Each method projects the data into a two-dimensional latent space, on which I visualized the results and quantitatively evaluated their ability to separate the ten digit classes.

The results of this experiment are summarized below.

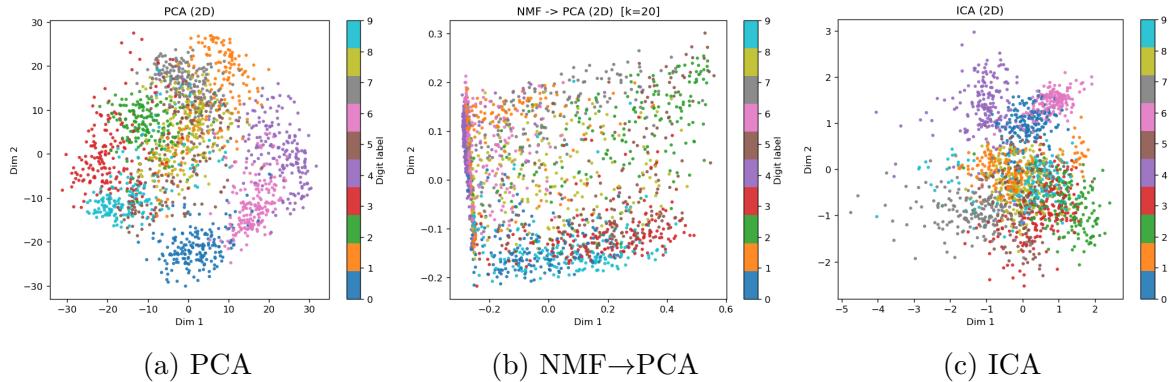


Figure 1: 2D embeddings of the digits data using PCA, NMF→PCA, and ICA. Colors denote true digit labels.

Method	ARI	NMI	Silhouette
PCA	0.3614	0.5190	0.3993
NMF → PCA	0.1588	0.2961	0.4080
ICA	0.3175	0.4567	0.3673

Table 1: Quantitative comparison of linear dimension reduction methods. The best scores for each metric are bolded.

PCA. For PCA, I retained the first two principal components and projected the samples into the 2D subspace they span, which preserves the primary directions of variance. As for hyperparameters, PCA has very few tunable parameters—the main one being the number of principal components. For ease of visualization, I set the number of components to 2

(PC=2) in this experiment. Figure 1a shows that the data are well dispersed, and digits such as 0, 3, 4, 6, and 9 form clearly separated clusters. Quantitatively, PCA achieved an ARI of 0.3614, an NMI of 0.5190, and a Silhouette score of 0.3993. Except for the Silhouette score, PCA obtained the highest values among the three methods. This indicates that PCA effectively separates the digits and maintains a high level of consistency with the true labels. Although its Silhouette value (approximately 0.4) is not the highest, it still suggests reasonably compact and well-separated clusters. This minor difference can be attributed to slight overlaps between neighboring clusters in the 2D embedding, even though the overall structure aligns well with the ground-truth classes.

In addition to the 2D embedding, I plotted the PCA scree plot and a bar chart of the top ten principal components' explained variance, as shown in Figure 2. Both plots reveal that the first three components contribute substantially more variance than the rest, with the third component explaining slightly less variance than the first two but significantly more than the fourth. This supports the observation that the 2D projection loses some discriminative information, which explains why the best ARI achieved by PCA remains moderate (0.3614).

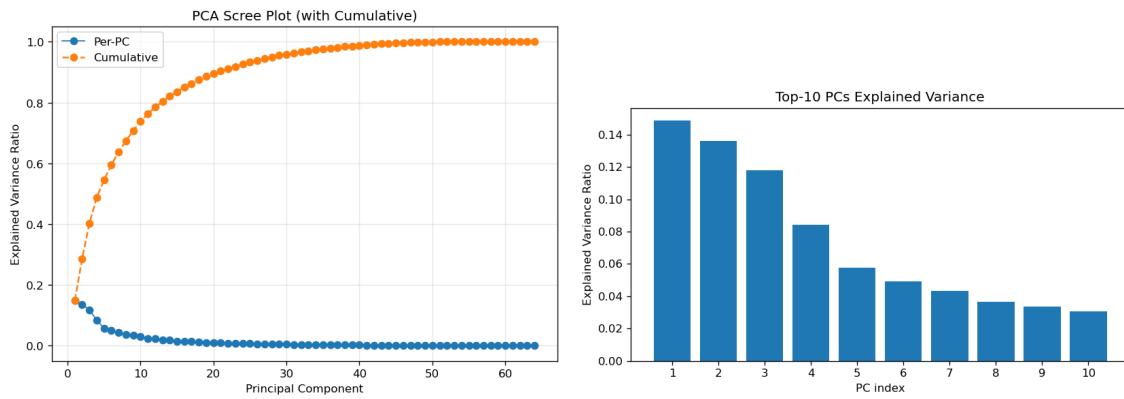


Figure 2: (Left) Scree plot showing the variance explained by each principal component; (Right) bar chart of the top-10 PCs' explained variance ratios.

Furthermore, I visualized the top ten PCA component images (Figure 3), which illustrate the principal modes of variation across digits.

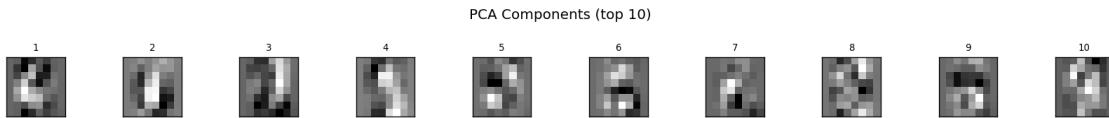


Figure 3: Top ten PCA components visualized as 8×8 basis images.

NMF→PCA. For the NMF method, I first determined the optimal number of components by minimizing the reconstruction error, which yielded $k = 20$. This means the data were decomposed into 20 non-negative basis vectors. The resulting coefficient matrix W was

then compressed into a 2D embedding space using PCA for visualization and comparison purposes, and the final result is shown in Figure 1(b).

As observed in the plot, NMF fails to clearly separate the digits in the 2D space. This is consistent with the quantitative results, where NMF achieved the lowest ARI (0.1588) and NMI (0.2961) among all methods. These findings indicate that the NMF representation captures local, part-based features rather than global discriminative structures, and that the subsequent PCA compression may distort these part-based patterns, reducing the overall interpretability.

On the other hand, NMF obtained the highest Silhouette score (0.4080), slightly higher than PCA. The embedding shows that nearly all samples are concentrated in the positive subspace, with only a small portion extending below zero (no less than -0.2). Even after PCA compression, this non-negativity-induced structure is largely preserved, resulting in an asymmetric distribution that nevertheless exhibits slightly better cluster compactness than PCA. This suggests that some of the features extracted by NMF may be better suited to local clustering in this dataset. In particular, compared to PCA, NMF tends to focus on localized regions of variation rather than global variance directions, which likely contributes to the formation of tighter, more compact clusters.

Similar to the PCA case where excluding the third principal component led to a loss of explanatory power, it is plausible that the 2D projection of NMF also omits important structural information. Specifically, the visualization reveals that the first principal component successfully separates digit “4” along the left margin, while the second principal component distinguishes digits “9” (light blue) and “0” from the rest of the samples, forming a clear boundary near the bottom region of the plot. This indicates that the first PC primarily captures the unique pattern of the digit “4”, whereas the second PC isolates the shapes shared by “0” and “9”. A potential 3D embedding including an additional principal component might further clarify the remaining overlapping clusters visible in the upper-right portion of the 2D space.

Finally, I visualized the NMF basis components, as shown in Figure 4. These components represent localized stroke-like patterns, providing an interpretable decomposition of the digits into additive parts.

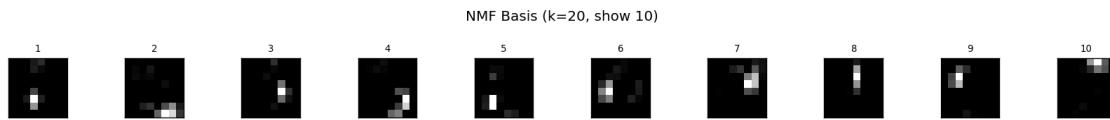


Figure 4: Top ten PCA components visualized as 8×8 basis images.

ICA. For the ICA method, I applied FastICA with two components after standardizing the data to ensure consistent feature scaling. As shown in Figure 1(c) and the summary table, ICA exhibits the most balanced overall performance among the three linear methods. Some clusters—notably digits 4, 6, and 0—are well separated and relatively compact, outperforming both PCA and NMF in local structure preservation. However, the remaining digits appear more mixed than in PCA, though less dispersed than in NMF. Visually,

the embedding forms two major groups: one consisting primarily of digits 4, 6, and 0, and another containing the other seven digits.

Quantitatively, ICA's ARI (0.3175) and NMI (0.4567) values lie between those of PCA and NMF, while its Silhouette score (0.3673) is the lowest among the three. This aligns with the visual interpretation: ICA achieves moderate global separability but weaker overall cluster compactness.

Regarding hyperparameters, ICA provides limited tuning options. Here I set `components=2`, reducing the data to a 2D space, and standardized all features prior to decomposition. Standardization is a conventional preprocessing step for ICA, as its underlying assumption relies on statistical independence between components. Consequently, the resulting embedding appears nearly spherical and centered around the origin—a distribution consistent with ICA's model assumptions but inconsistent with the inherent non-spherical structure of handwritten digits. This mismatch explains the weaker performance of ICA in this task.

Finally, the ICA component images (Figure 5) reveal contrast-like and edge-detecting patterns that emphasize stroke boundaries, differing from the smoother, global variations captured by PCA and the localized parts extracted by NMF.

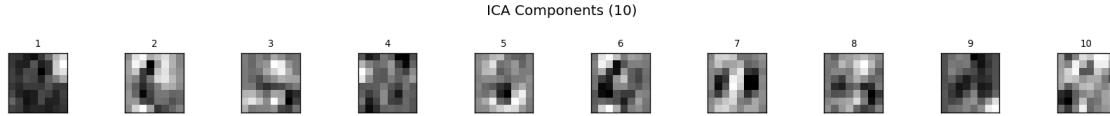


Figure 5: Top ten PCA components visualized as 8×8 basis images.

Overall Discussion. In summary, among the three linear dimension reduction techniques, PCA stands out as the most effective approach. It achieved the highest ARI (0.3614) and NMI (0.5190), and a Silhouette score (0.3993) close to the best. Visually, PCA produced the clearest and most interpretable clusters in the 2D embedding, separating several digits (such as 0, 3, 4, 6, and 9) distinctly. Given that PCA captures global variance directions, its performance would likely improve further in higher-dimensional embeddings, where additional components could better preserve discriminative variance.

1.2 Apply manifold learning approaches.

Methodology Each model was configured with typical hyperparameters and applied to the same standardized input features. The detailed settings are summarized below:

- **Kernel PCA (RBF)** — RBF kernel with $\gamma = 0.102$, chosen automatically based on the median pairwise distance.
- **Spectral Embedding** — Laplacian eigenmap-based manifold embedding with $n_{\text{neighbors}} = 10$.
- **Classical MDS** — Closed-form double-centered Euclidean distance eigendecomposition.

- **Metric MDS** — Stress minimization-based metric scaling via gradient descent.
- **t-SNE** — Perplexity of 30, PCA initialization, adaptive learning rate.
- **UMAP** — $n_{neighbors} = 15$, $\text{min_dist} = 0.1$.
- **Autoencoder (PyTorch)** — Fully connected symmetric architecture with input 64, hidden layer 32, bottleneck dimension 2, ReLU activations, sigmoid reconstruction, trained for 50 epochs with batch size 128 and learning rate 10^{-3} .

All embeddings were evaluated in 2D using k -means clustering ($k = 10$) with three quantitative metrics: Adjusted Rand Index (ARI), Normalized Mutual Information (NMI), and Silhouette coefficient.

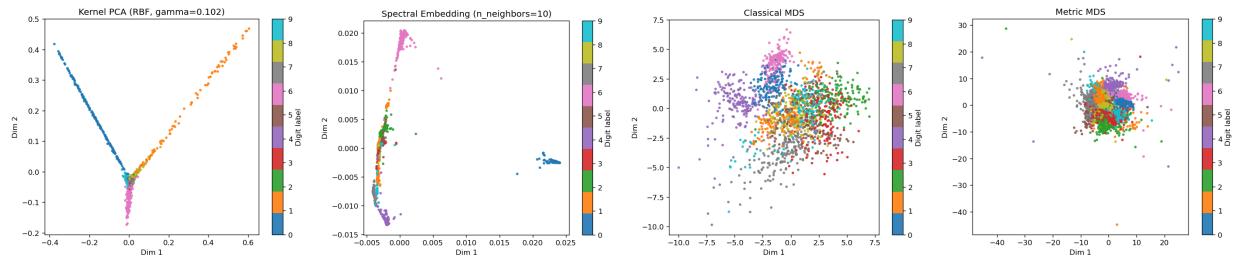


Figure 6: 2D embeddings of the digits dataset using kernel PCA, spectral embedding, classical mds and metric mds.

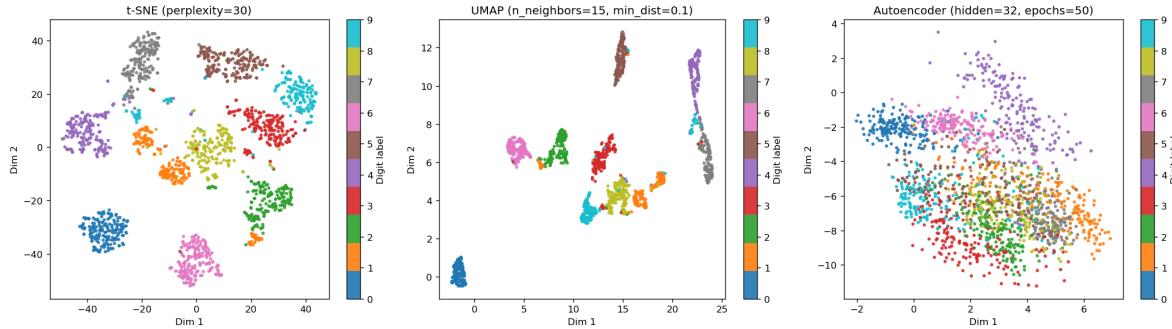


Figure 7: 2D embeddings of the digits dataset using t-SNE, UMAP and Pytorch Autoencoder.

Discussion The quantitative results and visualizations reveal clear distinctions among methods. UMAP and t-SNE perform best overall, yielding the strongest alignment with true digit labels—UMAP achieves the highest ARI (0.866) and NMI (0.891), while t-SNE produces similarly well-separated and visually coherent clusters. Spectral Embedding captures both global and local structures but exhibits partial overlap between classes. Kernel PCA attains the highest Silhouette score (0.754), reflecting locally compact but label-misaligned clusters, whereas Classical and Metric MDS show limited nonlinear capacity. The autoencoder provides a moderate nonlinear representation (ARI 0.35, NMI 0.50), illustrating

Method	ARI	NMI	Silhouette
Kernel PCA	0.0511	0.3357	0.7541
Spectral Embedding	0.4730	0.6681	0.6119
Classical MDS	0.3262	0.4647	0.3771
Metric MDS	0.3461	0.4722	0.3503
t-SNE	0.7695	0.8339	0.5731
UMAP	0.8662	0.8912	0.6866
Autoencoder	0.3476	0.4992	0.3958

Table 2: Quantitative comparison of seven embedding methods on the digits dataset ($k = 10$).

its potential but also the constraint of a shallow 2D bottleneck. Overall, UMAP and t-SNE best uncover the intrinsic manifold structure of handwritten digits, highlighting the strength of nonlinear embedding methods in revealing meaningful low-dimensional organization in high-dimensional data.

1.3 Discussion

Nonlinear manifold methods such as **UMAP** and **t-SNE** outperform linear approaches because they better preserve both local neighborhood relationships and global manifold geometry, thus capturing the intrinsic nonlinear structure of handwritten digits. Among all methods, **UMAP** performs best, achieving the highest ARI (0.866) and NMI (0.891) while forming compact and well-separated clusters that closely align with true digit classes. The most representative visualization is the **UMAP 2D embedding**, which clearly displays ten distinct clusters corresponding to the ten digits.

2 Open-Ended Data Analysis - Breast Cancer gene expression data.

2.1 Preprocessing

After loading the data (shape 445×359), six clinical columns were identified: `subtype`, `er_status`, `pr_status`, `her2_status`, `node`, and `metastasis`, leaving 353 gene expression features. No missing values or zero-variance genes were found, indicating high data quality. Each gene feature was standardized via Z-score normalization. Descriptive statistics of the clinical variables are summarized below:

- **Subtype:** Luminal A (200), Luminal B (106), Basal-like (79), HER2-enriched (53), Normal-like (7).
- **ER-Status:** Positive (339), Negative (100), Performed but Not Available (2), Indeterminate (2), Not Performed (2).

- **PR-Status:** Positive (291), Negative (147), Indeterminate (3), Performed but Not Available (2), Not Performed (2).
- **HER2-Status:** Negative (371), Positive (65), Equivocal (5), Not Available (4).
- **Node:** mean = 0.73, std = 0.87, range = [0, 3].
- **Metastasis:** mean = 0.025, std = 0.155 (mostly non-metastatic samples).

2.2 Methodology

Five dimension reduction methods were applied to the standardized gene expression matrix to obtain two-dimensional embeddings:

1. **Principal Component Analysis (PCA)** — linear orthogonal projection capturing maximal variance.
2. **Non-negative Matrix Factorization (NMF)** — parts-based representation using non-negative constraints (run on non-standardized data to ensure non-negativity).
3. **Spectral Embedding** — manifold learning based on graph Laplacian eigenvectors.
4. **t-SNE** — nonlinear embedding preserving local neighborhood structures.
5. **UMAP** — manifold approximation balancing local and global relationships.

Each embedding was visualized by coloring the points according to all six available clinical variables. However, for brevity and visual clarity, only the results colored by **Subtype** are shown here as representative examples. This selection illustrates overall trends while the complete set of figures (for all six variables across all five methods) is provided in the supplementary materials.

To quantitatively evaluate clustering quality with respect to molecular subtypes, k -means clustering ($k = 5$) was applied to each embedding, and three metrics were computed: Adjusted Rand Index (ARI), Normalized Mutual Information (NMI), and Silhouette coefficient.

2.3 Results

Overall, nonlinear manifold-based methods (t-SNE and UMAP) outperformed linear approaches in terms of ARI and NMI, suggesting that the underlying structure of gene expression data is highly nonlinear. UMAP achieved the highest Silhouette score (0.44), indicating that it produces well-separated clusters with clear boundaries, while t-SNE yielded the best ARI (0.21) and NMI (0.27), showing the strongest alignment with the known PAM50 subtypes. PCA achieved moderate performance, confirming its ability to preserve global variance but limited capacity to capture complex nonlinear relations. NMF performed comparably, producing slightly denser local clusters (reflected in its higher Silhouette score) but weaker subtype separation. Spectral embedding produced the most compact clusters but with limited biological interpretability due to its sensitivity to graph construction.

Method	ARI	NMI	Silhouette
PCA	0.1857	0.2399	0.3311
NMF	0.1712	0.2536	0.3842
Spectral	0.1421	0.2300	0.4020
t-SNE	0.2074	0.2704	0.3984
UMAP	0.2028	0.2780	0.4376

Table 3: Quantitative comparison of five dimension reduction methods on the BRCA dataset. Evaluation is based on clustering alignment with PAM50 subtypes ($k = 5$).

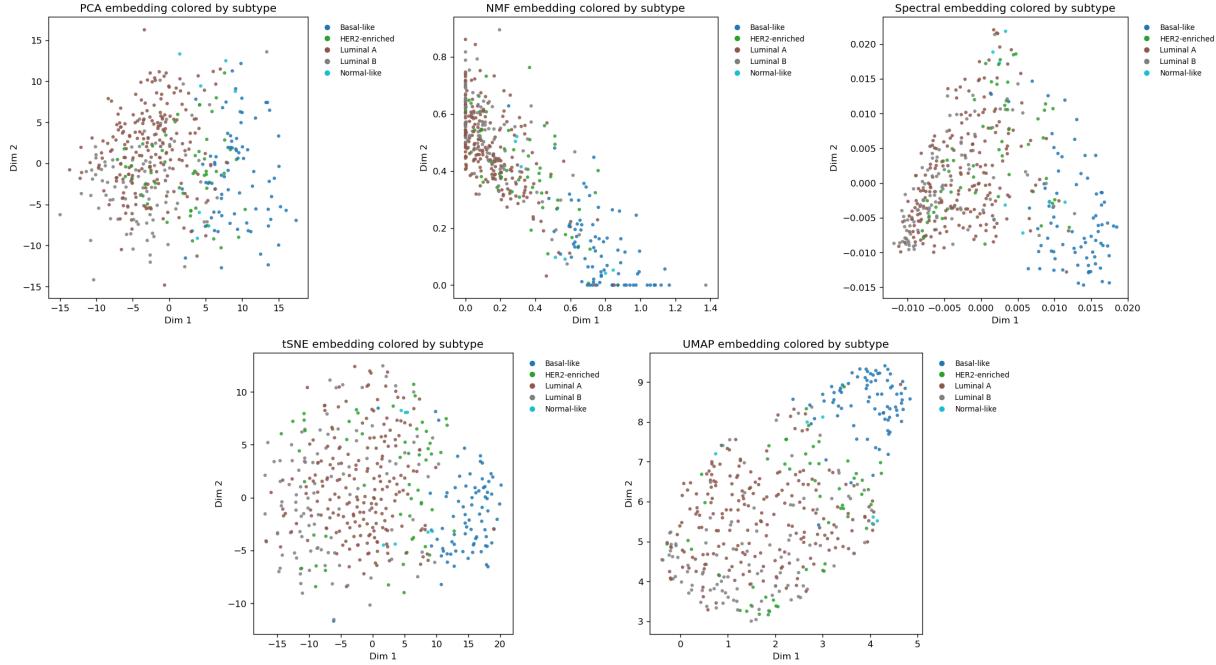


Figure 8: Two-dimensional embeddings of the BRCA gene expression data using five methods (PCA, NMF, Spectral, t-SNE, and UMAP), colored by molecular Subtype. All embeddings were also generated for the remaining five clinical variables, but only Subtype-colored results are displayed here for clarity.

2.4 Discussion

Given that no missing or low-variance genes were present, preprocessing had minimal impact on the raw data distribution. The observed performance differences mainly stem from the intrinsic characteristics of each method. The superior performance of UMAP and t-SNE suggests that manifold learning effectively captures nonlinear relationships among gene expressions that correspond to known molecular subtypes. Future work could extend this analysis by increasing latent dimensionality or incorporating autoencoders for deeper non-linear representations.

A Appendix: Code Implementation

A.1 Problem 1a.

```

1 """
2 Author: Chuyang Su cs4570@columbia.edu
3 Date: 2025-10-08 17:16:54
4 LastEditTime: 2025-10-08 18:21:50
5 FilePath: /Unsupervised-Learning-Homework/Homework 1/Code/Problem_1_a.py
6 Description:
7     This is the code part of GR5244 Unsupervised Learning Homework 1 Part 1a.
8 """
9 import argparse
10 import os
11 from pathlib import Path
12
13 import numpy as np
14 import matplotlib.pyplot as plt
15
16 from sklearn.datasets import load_digits
17 from sklearn.decomposition import PCA, NMF, FastICA
18 from sklearn.cluster import KMeans
19 from sklearn.metrics import adjusted_rand_score, normalized_mutual_info_score,
20                         silhouette_score
21 from sklearn.preprocessing import StandardScaler
22 from sklearn.pipeline import make_pipeline
23 import csv
24
25 # -----
26 # Utils
27 # -----
28 def ensure_dir(p):
29     Path(p).mkdir(parents=True, exist_ok=True)
30
31 def plot_embedding(X2d, y, title, outpath):
32     plt.figure(figsize=(6, 5), dpi=120)
33     scatter = plt.scatter(X2d[:, 0], X2d[:, 1], c=y, s=12, cmap='tab10', alpha
34                         =0.9, edgecolors='none')
35     plt.title(title)
36     plt.xlabel("Dim 1")
37     plt.ylabel("Dim 2")
38     cbar = plt.colorbar(scatter, ticks=range(10))
39     cbar.set_label("Digit label")
40     plt.tight_layout()
41     plt.savefig(outpath, bbox_inches='tight')
42     plt.close()
43
44 def plot_components(components, img_shape, n_show, title, outpath):
45     n = min(n_show, components.shape[0])
46     cols = 10 if n >= 10 else n
47     rows = int(np.ceil(n / cols))

```

```
48 plt.figure(figsize=(1.4*cols, 1.4*rows), dpi=120)
49 for i in range(n):
50     ax = plt.subplot(rows, cols, i + 1)
51     ax.imshow(components[i].reshape(img_shape), cmap='gray')
52     ax.set_xticks([])
53     ax.set_yticks([])
54     ax.set_title(f"{i+1}", fontsize=8)
55 plt.suptitle(title, y=1.02)
56 plt.tight_layout()
57 plt.savefig(outpath, bbox_inches='tight')
58 plt.close()

59
60 def plot_pca_scree(explained_var_ratio, outpath):
61     """Scree plot: per-PC variance ratio + cumulative curve"""
62     import numpy as np
63     import matplotlib.pyplot as plt
64
65     r = np.array(explained_var_ratio)
66     cum = np.cumsum(r)
67
68     plt.figure(figsize=(7, 5), dpi=120)
69     #
70     plt.plot(np.arange(1, len(r)+1), r, marker='o', linewidth=1)
71     #
72     plt.plot(np.arange(1, len(r)+1), cum, marker='o', linestyle='--')
73     plt.xlabel("Principal Component")
74     plt.ylabel("Explained Variance Ratio")
75     plt.title("PCA Scree Plot (with Cumulative)")
76     plt.legend(["Per-PC", "Cumulative"])
77     plt.grid(alpha=0.3)
78     plt.tight_layout()
79     plt.savefig(str(outpath), bbox_inches='tight')
80     plt.close()

81
82
83 def plot_pca_topk_bar(explained_var_ratio, k, outpath):
84     """Bar chart for top-k PCs' explained variance ratio"""
85     import numpy as np
86     import matplotlib.pyplot as plt
87
88     r = np.array(explained_var_ratio)[:k]
89     idx = np.arange(1, len(r)+1)

90
91     plt.figure(figsize=(7, 4), dpi=120)
92     plt.bar(idx, r)
93     plt.xlabel("PC index")
94     plt.ylabel("Explained Variance Ratio")
95     plt.title(f"Top-{k} PCs Explained Variance")
96     plt.xticks(idx)
97     plt.tight_layout()
98     plt.savefig(str(outpath), bbox_inches='tight')
99     plt.close()

100
101
```

```

102
103 def kmeans_scores(X2d, y, seed):
104     km = KMeans(n_clusters=10, n_init='auto', random_state=seed)
105     labels = km.fit_predict(X2d)
106     ari = adjusted_rand_score(y, labels)
107     nmi = normalized_mutual_info_score(y, labels)
108     sil = silhouette_score(X2d, labels)
109     return ari, nmi, sil
110
111
112 # -----
113 # Main pipeline for 1(a)
114 # -----
115 def run(seed=0, nmf_ks=(10, 15, 20), outdir="Homework 1/Latex"):
116     ensure_dir(outdir)
117
118     # Load data
119     digits = load_digits()
120     X = digits.data.astype(float)
121     y = digits.target
122     img_shape = (8, 8)
123
124     # -----
125     # PCA (2D embedding + components)
126     # -----
127     pca_2 = PCA(n_components=2, random_state=seed)
128     X_pca_2 = pca_2.fit_transform(X)
129     plot_embedding(X_pca_2, y, "PCA (2D)", f"{outdir}/pca_2d.png")
130
131     # For components visualization, use more PCs (e.g., 10)
132     pca_10 = PCA(n_components=10, random_state=seed).fit(X)
133     plot_components(pca_10.components_, img_shape, n_show=10,
134                      title="PCA Components (top 10)", outpath=f"{outdir}/
135                           pca_components.png")
136
137     pca_ari, pca_nmi, pca_sil = kmeans_scores(X_pca_2, y, seed)
138
139     pca_full = PCA(n_components=min(X.shape), random_state=seed).fit(X)
140     explained = pca_full.explained_variance_ratio_
141
142     plot_pca_scree(explained, Path(outdir) / "pca_scree.png")
143     plot_pca_topk_bar(explained, k=10, outpath=Path(outdir) / "pca_top10_var.
144                               png")
145
146     # -----
147     # NMF (grid over k), then 2D via PCA-on-W for visualization
148     # -----
149     X_nonneg = X - X.min() if X.min() < 0 else X
150     best_nmf = None
151     best_rec = np.inf
152     best_k = None

```

```

153     for k in nmf_ks:
154         nmf = NMF(n_components=k, init='nndsvda', random_state=seed, max_iter=
155                     1000)
156         W = nmf.fit_transform(X_nonneg)
157         rec = nmf.reconstruction_err_
158         if rec < best_rec:
159             best_rec = rec
160             best_nmf = nmf
161             best_k = k
162
163     # Use the best NMF
164     W_best = best_nmf.transform(X_nonneg) # (n_samples, best_k)
165     # Reduce W to 2D for visualization
166     nmf_to2 = PCA(n_components=2, random_state=seed)
167     X_nmf_2 = nmf_to2.fit_transform(W_best)
168     plot_embedding(X_nmf_2, y, f"NMF -> PCA (2D) [k={best_k}]", f"{outdir}/
169     nmf_2d.png")
170
171     # Visualize NMF basis (H)
172     plot_components(best_nmf.components_, img_shape, n_show=10,
173                     title=f"NMF Basis (k={best_k}, show 10)", outpath=f"{outdir}/
174                     nmf_components.png")
175
176     # -----
177     # ICA (2D embedding + components)
178     # -----
179     ica_pipeline = make_pipeline(StandardScaler(with_std=True), FastICA(
180                                     n_components=2, random_state=seed,
181                                     max_iter=1000))
182     X_ica_2 = ica_pipeline.fit_transform(X)
183     plot_embedding(X_ica_2, y, "ICA (2D)", f"{outdir}/ica_2d.png")
184
185     # For components display, fit a separate ICA with more comps (e.g., 10) on
186     # standardized X
187     scaler = StandardScaler(with_std=True)
188     X_std = scaler.fit_transform(X)
189     ica_10 = FastICA(n_components=10, random_state=seed, max_iter=1000).fit(
190                     X_std)
191     plot_components(ica_10.mixing_.T, img_shape, n_show=10, # mixing_.T
192                     component ""images
193                     title="ICA Components (10)", outpath=f"{outdir}/
194                     ica_components.png")
195
196     ica_ari, ica_nmi, ica_sil = kmeans_scores(X_ica_2, y, seed)
197
198     # -----
199     # Save metrics
200     # -----
201     metrics_path = f"{outdir}/part1a_metrics.csv"

```

```

195     with open(metrics_path, "w", newline="") as f:
196         writer = csv.writer(f)
197         writer.writerow(["Method", "Params", "ARI", "NMI", "Silhouette", "Notes"])
198         writer.writerow(["PCA", "n_components=2", f"{pca_ari:.4f}", f"{pca_nmi:.4f}", f"{pca_sil:.4f}", "2D embedding"])
199         writer.writerow(["NMF -> PCA", f"k={best_k}, then 2D PCA", f"{nmf_ari:.4f}", f"{nmf_nmi:.4f}", f"{nmf_sil:.4f}",
200                         f"best reconstruction k among {list(nmf_ks)} (err={best_rec:.4f})"])
201         writer.writerow(["ICA", "n_components=2 (with standardization)", f"{ica_ari:.4f}", f"{ica_nmi:.4f}",
202                         f"{ica_sil:.4f}", "2D embedding"])
203
204     # Also print a neat summary
205     print("\n==== Part 1(a) - Linear Methods on Digits ===")
206     print(f"[PCA] ARI={pca_ari:.4f} NMI={pca_nmi:.4f} Silhouette={pca_sil:.4f}")
207     print(f"[NMF -> PCA] ARI={nmf_ari:.4f} NMI={nmf_nmi:.4f} Silhouette={nmf_sil:.4f} (best k={best_k},
208           recon_err={best_rec:.4f})")
209     print(f"[ICA] ARI={ica_ari:.4f} NMI={ica_nmi:.4f} Silhouette={ica_sil:.4f}")
210
211     print(f"\nSaved figures and metrics to: {outdir}/")
212     print("Figures:")
213     print(" - pca_2d.png, pca_components.png")
214     print(" - nmf_2d.png, nmf_components.png")
215     print(" - ica_2d.png, ica_components.png")
216     print("Table:")
217     print(f" - {Path(metrics_path).name}")
218
219 # -----
220 # Entry
221 # -----
222 if __name__ == "__main__":
223     parser = argparse.ArgumentParser()
224     parser.add_argument("--seed", type=int, default=25)
225     parser.add_argument("--outdir", type=str, default="Homework 1/Latex/Results/Problem_1_a")
226     parser.add_argument("--nmf_ks", type=int, nargs="+", default=[10, 15, 20])
227     args = parser.parse_args()
228     run(seed=args.seed, nmf_ks=tuple(args.nmf_ks), outdir=args.outdir)

```

A.2 Problem 1b.

```

1 #!/usr/bin/env python3
2 # -*- coding: utf-8 -*-

```

```
3  """
4  Author: Chuyang Su <cs4570@columbia.edu>
5  Date: 2025-10-10
6  FilePath: /Unsupervised-Learning-Homework/Homework 1/Code/Problem_1_b.py
7  Description:
8      GR5244 HW1 Part 1(b): Manifold learning on sklearn digits (n=1797, p=64).
9          Methods: Kernel PCA, Spectral Embedding, Classical MDS, Metric MDS, t-
10             SNE, UMAP, Autoencoder.
11
12 from pathlib import Path
13 import argparse
14 import numpy as np
15 import matplotlib.pyplot as plt
16
17 from sklearn.datasets import load_digits
18 from sklearn.preprocessing import StandardScaler, MinMaxScaler
19 from sklearn.decomposition import KernelPCA
20 from sklearn.manifold import SpectralEmbedding, MDS, TSNE
21 from sklearn.metrics import adjusted_rand_score, normalized_mutual_info_score,
22                             silhouette_score, pairwise_distances
23 from sklearn.cluster import KMeans
24 import umap
25
26 import torch
27 import torch.nn as nn
28 import torch.optim as optim
29 from torch.utils.data import TensorDataset, DataLoader
30 # -----
31 # Utils
32 # -----
33 def ensure_dir(p):
34     Path(p).mkdir(parents=True, exist_ok=True)
35
36 def set_seed(seed: int):
37     import random, os
38     random.seed(seed); np.random.seed(seed)
39     torch.manual_seed(seed); torch.cuda.manual_seed_all(seed)
40     torch.backends.cudnn.deterministic = True
41     torch.backends.cudnn.benchmark = False
42
43 def get_device():
44     return torch.device("cuda" if torch.cuda.is_available() else "cpu")
45
46 def plot_embedding(X2d, y, title, outpath):
47     """Simple 2D scatter with digits colormap; no axes ticks; safe colormap
48     API."""
49     cmap = plt.colormaps.get_cmap('tab10')
50     plt.figure(figsize=(6, 5), dpi=120)
51     sc = plt.scatter(X2d[:, 0], X2d[:, 1], c=y, s=12, cmap=cmap, alpha=0.9,
52                      edgecolors='none')
53     plt.title(title)
54     plt.xlabel("Dim 1"); plt.ylabel("Dim 2")
```

```

53     cbar = plt.colorbar(sc, ticks=range(10))
54     cbar.set_label("Digit label")
55     plt.tight_layout()
56     plt.savefig(outpath, bbox_inches='tight')
57     plt.close()
58
59
60 def kmeans_scores(X2d, y, seed, k=10):
61     km = KMeans(n_clusters=k, random_state=seed, n_init=10)
62     labels = km.fit_predict(X2d)
63     ari = adjusted_rand_score(y, labels)
64     nmi = normalized_mutual_info_score(y, labels)
65     sil = silhouette_score(X2d, labels)
66     return ari, nmi, sil
67
68
69 def classical_mds(X, n_components=2):
70     """
71         Classical MDS (Torgerson/Gower): eigendecomposition of double-centered
72         squared-distance matrix. Returns low-d coords (can be rotated/reflected).
73     """
74     D = pairwise_distances(X, metric='euclidean')
75     D2 = D ** 2
76     n = D2.shape[0]
77     J = np.eye(n) - np.ones((n, n)) / n
78     B = -0.5 * J @ D2 @ J
79     # Eigen-decomposition
80     evals, evecs = np.linalg.eigh(B)
81     # Take top components
82     idx = np.argsort(evals)[::-1]
83     evals = evals[idx]; evecs = evecs[:, idx]
84     # Keep only positive eigenvalues
85     pos = evals > 0
86     evals = evals[pos]; evecs = evecs[:, pos]
87     evals_k = evals[:n_components]
88     evecs_k = evecs[:, :n_components]
89     X_emb = evecs_k * np.sqrt(np.maximum(evals_k, 0))
90     return X_emb
91
92
93 def auto_gamma_rbf(X):
94     """
95         Heuristic gamma for RBF kernel (Kernel PCA): gamma = 1 / median(pairwise
96         distance).
97         Scale-robust; avoids manual guesswork when not provided.
98     """
99     d = pairwise_distances(X, metric='euclidean')
100    med = np.median(d)
101    if med <= 0:
102        return 1.0
103    return 1.0 / med
104
105 # Autoencoder

```

```
106 class AE(nn.Module):
107     def __init__(self, input_dim: int, hidden: int = 32, bottleneck: int = 2):
108         super().__init__()
109         self.encoder = nn.Sequential(
110             nn.Linear(input_dim, hidden),
111             nn.ReLU(),
112             nn.Linear(hidden, bottleneck) #
113         )
114         self.decoder = nn.Sequential(
115             nn.Linear(bottleneck, hidden),
116             nn.ReLU(),
117             nn.Linear(hidden, input_dim),
118             nn.Sigmoid() # [0,1]
119     )
120
121     def forward(self, x):
122         z = self.encoder(x)
123         xhat = self.decoder(z)
124         return xhat, z
125
126
127 @torch.no_grad()
128 def encode_dataset(model: AE, loader: DataLoader, device):
129     model.eval()
130     zs = []
131     for (xb,) in loader:
132         xb = xb.to(device)
133         _, z = model(xb)
134         zs.append(z.cpu().numpy())
135     return np.concatenate(zs, axis=0)
136
137
138 def train_autoencoder_pytorch(
139     X_minmax: np.ndarray,
140     seed: int = 0,
141     hidden: int = 32,
142     bottleneck: int = 2,
143     epochs: int = 50,
144     batch_size: int = 128,
145     lr: float = 1e-3,
146 ):
147     set_seed(seed)
148     device = get_device()
149
150     X_tensor = torch.tensor(X_minmax, dtype=torch.float32)
151     ds = TensorDataset(X_tensor)
152     dl = DataLoader(ds, batch_size=batch_size, shuffle=True, drop_last=False)
153
154     model = AE(input_dim=X_minmax.shape[1], hidden=hidden, bottleneck=
155                 bottleneck).to(device)
156     opt = optim.Adam(model.parameters(), lr=lr)
157     crit = nn.MSELoss()
158     model.train()
159     for _ in range(epochs):
```

```

159     for (xb,) in dl:
160         xb = xb.to(device)
161         xhat, _ = model(xb)
162         loss = crit(xhat, xb)
163         opt.zero_grad()
164         loss.backward()
165         opt.step()
166
167     #      2D
168     #      DataLoader
169     Z = encode_dataset(model, DataLoader(ds, batch_size=1024, shuffle=False),
170                         device)
171
172
173 # -----
174 # Main pipeline for 1(b)
175 # -----
176 def run(seed=25,
177         outdir="Homework 1/Latex/Results/Problem_1_b",
178         kpca_gamma=None,
179         spectral_n_neighbors=10,
180         tsne_perplexity=30,
181         umap_n_neighbors=15,
182         umap_min_dist=0.1,
183         ae_hidden=32,
184         ae_epochs=50,
185         ae_batch=128,
186         ae_lr=1e-3):
187     set_seed(seed)
188     ensure_dir(outdir)
189
190     # Load data
191     digits = load_digits()
192     X = digits.data.astype(float)           # (1797, 64)
193     y = digits.target                      # (1797,)
194
195     # Standardize (common for manifold learning)
196     X_std = StandardScaler().fit_transform(X)
197
198     # 1) Kernel PCA (RBF)
199     gamma = kpca_gamma if kpca_gamma is not None else auto_gamma_rbf(X_std)
200     kpca = KernelPCA(n_components=2, kernel='rbf', gamma=gamma,
201                       fit_inverse_transform=False,
202                       random_state=seed)
203     X_kpca = kpca.fit_transform(X_std)
204     plot_embedding(X_kpca, y, f"Kernel PCA (RBF, gamma={gamma:.3g})", f"{outdir}/kpca_2d.png")
205     kpca_ari, kpca_nmi, kpca_sil = kmeans_scores(X_kpca, y, seed)
206
207     # 2) Spectral Embedding
208     se = SpectralEmbedding(n_components=2, n_neighbors=spectral_n_neighbors,
209                           random_state=seed)
210     X_se = se.fit_transform(X_std)

```

```

208 plot_embedding(X_se, y, f"Spectral Embedding (n_neighbors={spectral_n_neighbors})", f"{outdir}/spectral_2d.png")
209 se_ari, se_nmi, se_sil = kmeans_scores(X_se, y, seed)
210
211 # 3) Classical MDS (closed-form)
212 X_cmds = classical_mds(X_std, n_components=2)
213 plot_embedding(X_cmds, y, "Classical MDS", f"{outdir}/classical_mds_2d.png")
214 cmds_ari, cmds_nmi, cmds_sil = kmeans_scores(X_cmds, y, seed)
215
216 # 4) Metric MDS (sklearn MDS, metric=True)
217 mds_metric = MDS(n_components=2, metric=True, normalized_stress='auto',
218                   random_state=seed)
219 X_mds = mds_metric.fit_transform(X_std)
220 plot_embedding(X_mds, y, "Metric MDS", f"{outdir}/metric_mds_2d.png")
221 mds_ari, mds_nmi, mds_sil = kmeans_scores(X_mds, y, seed)
222
223 # 5) t-SNE
224 tsne = TSNE(n_components=2, perplexity=tsne_perplexity, learning_rate='auto',
225               init='pca', random_state=seed)
226 X_tsne = tsne.fit_transform(X_std)
227 plot_embedding(X_tsne, y, f"t-SNE (perplexity={tsne_perplexity})", f"{outdir}/tsne_2d.png")
228 tsne_ari, tsne_nmi, tsne_sil = kmeans_scores(X_tsne, y, seed)
229
230 # 6) UMAP
231 umap_model = umap.UMAP(n_neighbors=umap_n_neighbors, min_dist=umap_min_dist,
232                         n_components=2, random_state=seed)
233 X_umap = umap_model.fit_transform(X_std)
234 plot_embedding(X_umap, y, f"UMAP (n_neighbors={umap_n_neighbors}, min_dist={umap_min_dist})", f"{outdir}/umap_2d.png")
235 umap_ari, umap_nmi, umap_sil = kmeans_scores(X_umap, y, seed)
236
237 # AutoencoderPyTorch
238 X_minmax = MinMaxScaler().fit_transform(X)
239 X_ae = train_autoencoder_pytorch(
240     X_minmax,
241     seed=seed,
242     hidden=ae_hidden,
243     bottleneck=2,
244     epochs=ae_epochs,
245     batch_size=ae_batch,
246     lr=ae_lr,
247 )
248 plot_embedding(X_ae, y, f"Autoencoder (hidden={ae_hidden}, epochs={ae_epochs})", f"{outdir}/autoencoder_2d.png")
249 ae_ari, ae_nmi, ae_sil = kmeans_scores(X_ae, y, seed)
250
251 # Save metrics

```

```

249 import csv
250 metrics_path = f"{outdir}/part1b_metrics.csv"
251 with open(metrics_path, "w", newline="") as f:
252     writer = csv.writer(f)
253     writer.writerow(["Method", "Params", "ARI", "NMI", "Silhouette", "Notes"])
254     writer.writerow(["Kernel PCA", f"rbf", gamma={gamma:.3g}, f"{kpca_ari:.4f}", f"{kpca_nmi:.4f}", f"{kpca_sil:.4f}", "2D embedding"])
255     writer.writerow(["Spectral", f"n_neighbors={spectral_n_neighbors}", f"{se_ari:.4f}", f"{se_nmi:.4f}", f"{se_sil:.4f}", "2D embedding"])
256     writer.writerow(["Classical MDS", "-", f"{cmds_ari:.4f}", f"{cmds_nmi:.4f}", f"{cmds_sil:.4f}", "closed-form (eigendecomposition")])
257     writer.writerow(["Metric MDS", "sklearn MDS(metric=True)", f"{mds_ari:.4f}", f"{mds_nmi:.4f}", f"{mds_sil:.4f}", "stress minimization"])
258     writer.writerow(["t-SNE", f"perplexity={tsne_perplexity}", f"{tsne_ari:.4f}", f"{tsne_nmi:.4f}", f"{tsne_sil:.4f}", "2D embedding"])
259     writer.writerow(["UMAP", f"n_neighbors={umap_n_neighbors}", min_dist={umap_min_dist}, f"{umap_ari:.4f}", f"{umap_nmi:.4f}", f"{umap_sil:.4f}", "2D embedding"])
260     writer.writerow(["Autoencoder", f"hidden={ae_hidden}", epochs={ae_epochs}, batch={ae_batch}, lr={ae_lr:g}, f"{ae_ari:.4f}", f"{ae_nmi:.4f}", f"{ae_sil:.4f}", "2D bottleneck (PyTorch)"])
261
262 # Print a neat summary
263 print("\n== Part 1(b) - Manifold Methods on Digits ==")
264 print(f"[Kernel PCA] ARI={kpca_ari:.4f} NMI={kpca_nmi:.4f} Silhouette={kpca_sil:.4f} (gamma={gamma:.3g})")
265 print(f"[Spectral] ARI={se_ari:.4f} NMI={se_nmi:.4f} Silhouette={se_sil:.4f} (n_neighbors={spectral_n_neighbors})")
266 print(f"[Classical MDS] ARI={cmds_ari:.4f} NMI={cmds_nmi:.4f} Silhouette={cmdssil:.4f}")
267 print(f"[Metric MDS] ARI={mds_ari:.4f} NMI={mds_nmi:.4f} Silhouette={mds_sil:.4f}")
268 print(f"[t-SNE] ARI={tsne_ari:.4f} NMI={tsne_nmi:.4f} Silhouette={tsne_sil:.4f} (perplexity={tsne_perplexity})")
269 print(f"[UMAP] ARI={umap_ari:.4f} NMI={umap_nmi:.4f} Silhouette={umap_sil:.4f} (n_neighbors={umap_nneighbors})")

```

```

                                umap_n_neighbors}, min_dist={
                                umap_min_dist})")
270 print(f"\nSaved figures and metrics to: {outdir}/")
271 print("Figures:")
272 print(" - kPCA_2d.png, spectral_2d.png, classical_mds_2d.png,
273                                metric_mds_2d.png, tSNE_2d.png,
274                                UMAP_2d.png")
275 print("Table:")
276 print(f" - {Path(metrics_path).name}")
277 print(f"[Autoencoder] ARI={ae_ari:.4f} NMI={ae_nmi:.4f} Silhouette={ae_sil:.4f} (hidden={ae_hidden}, epochs={ae_epochs})")
278
279 # -----
280 # Entry
281 # -----
282
283 if __name__ == "__main__":
284     parser = argparse.ArgumentParser()
285     parser.add_argument("--seed", type=int, default=0, help="random seed")
286     parser.add_argument("--outdir", type=str, default="Homework 1/Latex/
287                               Results/Problem_1_b", help="output
288                               directory")
289     parser.add_argument("--kPCA_gamma", type=float, default=None, help="RBF
290                               gamma for Kernel PCA (auto if None)
291                               ")
292     parser.add_argument("--spectral_n_neighbors", type=int, default=10, help="n_neighbors for Spectral Embedding")
293     parser.add_argument("--tSNE_perplexity", type=float, default=30, help="perplexity for t-SNE")
294     parser.add_argument("--UMAP_n_neighbors", type=int, default=15, help="n_neighbors for UMAP")
295     parser.add_argument("--UMAP_min_dist", type=float, default=0.1, help="min_dist for UMAP")
296     parser.add_argument("--ae_hidden", type=int, default=32)
297     parser.add_argument("--ae_epochs", type=int, default=50)
298     parser.add_argument("--ae_batch", type=int, default=128)
299     parser.add_argument("--ae_lr", type=float, default=1e-3)
300     args = parser.parse_args()
301
302     run(seed=args.seed,
303         outdir=args.outdir,
304         kPCA_gamma=args.kPCA_gamma,
305         spectral_n_neighbors=args.spectral_n_neighbors,
306         tSNE_perplexity=args.tSNE_perplexity,
307         UMAP_n_neighbors=args.UMAP_n_neighbors,
308         UMAP_min_dist=args.UMAP_min_dist,
309         ae_hidden=args.ae_hidden,
310         ae_epochs=args.ae_epochs,
311         ae_batch=args.ae_batch,
312         ae_lr=args.ae_lr)

```

A.3 Problem 2.

```

43     elif method == 'tsne':
44         model = TSNE(n_components=n_components, random_state=seed, **kwargs)
45         name = "tSNE"
46     elif method == 'umap':
47         model = umap.UMAP(n_components=n_components, random_state=seed, **
48                             kwargs)
49         name = "UMAP"
50     elif method == 'spectral':
51         model = SpectralEmbedding(n_components=n_components, random_state=seed
52                                     )
53         name = "Spectral"
54     else:
55         raise ValueError(f"Unknown method: {method}")
56 X_embedded = model.fit_transform(X)
57 return X_embedded, name
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91 def evaluate_embedding(X2d, subtype_series, seed=0):

```

```

92     """KMeans(k=5) on the embedding; compute ARI/NMI/Silhouette using Subtype
93         as reference."""
94     mask = subtype_series.notna()
95     X_eval = X2d[mask.values]
96     y_ref = subtype_series[mask].astype(str).values
97
98     km = KMeans(n_clusters=5, random_state=seed, n_init=10)
99     pred = km.fit_predict(X_eval)
100
101    ari = adjusted_rand_score(y_ref, pred)
102    nmi = normalized_mutual_info_score(y_ref, pred)
103    sil = silhouette_score(X_eval, pred)
104    return ari, nmi, sil
105
106
107 def run_problem2_pipeline(
108     data_path="Homework 1/Data/BRCA_data.csv",
109     outdir="Homework 1/Latex/Results/Problem_2",
110     seed=0
111 ) :
112     outdir = Path(outdir)
113     outdir.mkdir(parents=True, exist_ok=True)
114
115     # 1) Preprocess (writes TXT report under outdir)
116     X, y_clinical, gene_cols = load_and_preprocess_brca(
117         data_path=data_path,
118         report_path=str(outdir / "preprocessing_report.txt")
119     )
120
121     # 2) Methods to run (UMAP included by assumption)
122     methods = [
123         ("pca",      dict()),
124         ("nmf",      dict()),
125         ("spectral", dict()),
126         ("tsne",     dict(perplexity=30, learning_rate='auto', init='pca')),
127         ("umap",     dict(n_neighbors=15, min_dist=0.1)),
128     ]
129
130     # 3) Clinical hues to color by (only those that exist)
131     hues = [c for c in ["subtype", "er_status", "pr_status", "her2_status"] if
132             c in y_clinical.columns]
133
134     # 4) Run, plot, evaluate
135     metrics_rows = []
136     best_by_ari = (None, -1.0)
137
138     for m, kwargs in methods:
139         if m == "nmf":
140             # NMF
141             X_nonneg = np.maximum(X, 0)  #
142             X2d, name = run_embedding(X_nonneg, method=m, seed=seed, **kwargs)
143         else:
144             X2d, name = run_embedding(X, method=m, seed=seed, **kwargs)

```

```

144     # Evaluation (Subtype as reference, if present)
145     if "subtype" in y_clinical.columns:
146         ari, nmi, sil = evaluate_embedding(X2d, y_clinical["subtype"],
147                                         seed=seed)
147     else:
148         ari = nmi = sil = np.nan
149
150     metrics_rows.append([name, f"{ari:.4f}", f"{nmi:.4f}", f"{sil:.4f}"])
151     if not np.isnan(ari) and ari > best_by_ari[1]:
152         best_by_ari = (name, ari)
153
154     # Save embedding as CSV (optional, useful for debugging/report)
155     emb_path = outdir / f"{name}_embedding.csv"
156     np.savetxt(emb_path, X2d, delimiter=",")
157
158     # Plots colored by clinical variables
159     for hue in hues:
160         fig_path = outdir / f"{name}_{hue}.png"
161         plot_embedding(X2d, y_clinical, hue_col=hue, title=name, outpath=
162                         fig_path)
162
163     # 5) Save metrics summary
164     metrics_path = outdir / "metrics_summary.csv"
165     with open(metrics_path, "w", newline="") as f:
166         writer = csv.writer(f)
167         writer.writerow(["Method", "ARI (Subtype)", "NMI (Subtype)", "
168                         Silhouette (KMeans=5)"])
168         writer.writerows(metrics_rows)
169
170     # Console summary
171     print("\n== Problem 2: embedding metrics (Subtype as reference) ===")
172     for r in metrics_rows:
173         print("{:,<10s} ARI={}  NMI={}  Sil={}" .format(*r))
174     if best_by_ari[0] is not None:
175         print(f"\nBest by ARI: {best_by_ari[0]} (ARI={best_by_ari[1]:.4f})")
176
177     print(f"\nSaved figures & tables to: {outdir}")
178
179 def load_and_preprocess_brca(data_path: str, var_threshold: float = 1e-4,
180                               report_path: str = None):
180     # 1) Load with sample IDs in index
181     df = pd.read_csv(data_path, index_col=0)
182     n_samples_raw, n_cols_raw = df.shape
183     print(f"Raw data shape: {n_samples_raw} samples x {n_cols_raw} columns")
184
185     # 2) Normalize column names
186     df.columns = (
187         df.columns.astype(str)
188         .str.strip()
189         .str.replace("-", "_", regex=False)
190         .str.replace(" ", "_", regex=False)
191         .str.lower()
192     )
193

```

```

194 # 3) Detect clinical columns robustly
195 possible_clinical = ["subtype", "er_status", "pr_status", "her2_status", " "
196   node", "metastasis"]
197 clinical_cols = [c for c in possible_clinical if c in df.columns]
198 if not clinical_cols:
199     raise ValueError("No clinical columns detected after normalization. "
200   "Please check the CSV headers.")
201 print(f"Detected clinical columns: {clinical_cols}")

202 # 4) Split gene vs clinical
203 gene_cols = [c for c in df.columns if c not in clinical_cols]
204 # Force numeric on genes; coerce errors to NaN
205 df[gene_cols] = df[gene_cols].apply(pd.to_numeric, errors="coerce")

206 X_df = df[gene_cols].copy()
207 y_clinical = df[clinical_cols].copy()

208 # 5) Missing value stats and imputation
209 n_missing_total = int(X_df.isna().sum().sum())
210 n_rows_with_na = int(X_df.isna().any(axis=1).sum())
211 n_genes_with_na = int(X_df.isna().any(axis=0).sum())
212 print(f"Missing values in gene matrix: {n_missing_total} "
213   f"(rows with any NA: {n_rows_with_na}, genes with any NA: {"
214   n_genes_with_na})")

215 # Impute clinical: categorical by mode, numeric by mean
216 for col in y_clinical.columns:
217     if y_clinical[col].dtype == "O":
218         mode_vals = y_clinical[col].mode(dropna=True)
219         fill_val = mode_vals.iloc[0] if not mode_vals.empty else "Unknown"
220         y_clinical[col] = y_clinical[col].fillna(fill_val)
221     else:
222         y_clinical[col] = y_clinical[col].fillna(y_clinical[col].mean())

223 # Impute genes by column median (robust)
224 X_df = X_df.fillna(X_df.median())

225 # 6) Standardize (Z-score)
226 scaler = StandardScaler()
227 X_scaled = scaler.fit_transform(X_df)

228 # 7) Build and optionally save summary report
229 summary_lines = [
230     "==== BRCA Data Preprocessing Summary ===",
231     f"Original shape: {n_samples_raw} samples x {n_cols_raw} columns",
232     f"Detected clinical columns: {clinical_cols}",
233     f"Initial gene columns: {len(gene_cols)}",
234     f"Total missing values filled (genes): {n_missing_total}",
235     f"Rows with any NA (genes): {n_rows_with_na}",
236     f"Genes with any NA: {n_genes_with_na}",
237     f"Final standardized data shape: {X_scaled.shape}",
238 ]
239
240 # --- Clinical summary by type ---

241
242
243
244
245

```

```

246     summary_text = ["Clinical variable summary:"]
247     for col in y_clinical.columns:
248         if y_clinical[col].dtype == "O":
249             counts = y_clinical[col].value_counts(dropna=False)
250             summary_text.append(f"\n{col} (categorical):")
251             summary_text.append(str(counts))
252         else:
253             desc = y_clinical[col].describe()
254             summary_text.append(f"\n{col} (numeric):")
255             summary_text.append(str(desc))
256
257     report_text = "\n".join(summary_lines + ["\n".join(summary_text)])
258     print("\n" + report_text)
259
260     if report_path is not None:
261         report_dir = Path(report_path).parent
262         report_dir.mkdir(parents=True, exist_ok=True)
263         with open(report_path, "w", encoding="utf-8") as f:
264             f.write(report_text)
265         print(f"\nPreprocessing summary saved to {report_path}")
266
267     return X_scaled, y_clinical, gene_cols
268
269 def run_embedding(X, method='pca', n_components=2, seed=0, **kwargs):
270     method = method.lower()
271     if method == 'pca':
272         model = PCA(n_components=n_components, random_state=seed)
273         name = 'PCA'
274     elif method == 'nmf':
275         model = NMF(n_components=n_components, init='nndsvda', random_state=
276                     seed, max_iter=1000)
276         name = 'NMF'
277     elif method == 'tsne':
278         model = TSNE(n_components=n_components, random_state=seed, **kwargs)
279         name = 'tSNE'
280     elif method == 'umap':
281         model = umap.UMAP(n_components=n_components, random_state=seed, **
282                            kwargs)
283         name = 'UMAP'
284     elif method == 'spectral':
285         model = SpectralEmbedding(n_components=n_components, random_state=seed
286                                   )
287         name = 'Spectral'
288     else:
289         raise ValueError(f"Unknown method: {method}")
290
291     X_emb = model.fit_transform(X)
292     return X_emb, name
293
294 def _encode_categories(series):
295     cats = series.astype(str).fillna("NA").values
296     uniq = sorted(np.unique(cats))
297     mapping = {c: i for i, c in enumerate(uniq)}

```

```

297     idx = np.array([mapping[c] for c in cats], dtype=int)
298     return idx, mapping
299
300
301 def plot_embedding(X2d, meta, hue_col, title, outpath):
302     outpath = Path(outpath)
303     outpath.parent.mkdir(parents=True, exist_ok=True)
304
305     labels = meta[hue_col].astype(str).fillna("NA")
306     idx, mapping = _encode_categories(labels)
307     cmap = plt.cm.get_cmap('tab10', len(mapping))
308
309     plt.figure(figsize=(6.2, 5.2), dpi=120)
310     sc = plt.scatter(X2d[:,0], X2d[:,1], c=idx, s=14, cmap=cmap, alpha=0.85,
311                      edgecolors='none')
312     plt.title(f"{title} embedding colored by {hue_col}")
313     plt.xlabel("Dim 1"); plt.ylabel("Dim 2")
314     handles = [plt.Line2D([0],[0], marker='o', linestyle='',
315                         markersize=6, markerfacecolor=cmap(i),
316                                         markeredgecolor
317                                         ='none')
318                         for i in range(len(mapping))]
319     labels_sorted = list(mapping.keys())
320     plt.legend(handles, labels_sorted, bbox_to_anchor=(1.02, 1), loc='upper
321                 left', fontsize=8, frameon=False)
322     plt.tight_layout()
323     plt.savefig(outpath, bbox_inches='tight')
324     plt.close()
325
326
327
328 def evaluate_embedding(X2d, subtype_series, seed=0):
329     mask = subtype_series.notna()
330     X_eval = X2d[mask.values]
331     y_ref = subtype_series[mask].astype(str).values
332
333     km = KMeans(n_clusters=5, random_state=seed, n_init=10)
334     pred = km.fit_predict(X_eval)
335
336
337
338
339 if __name__ == "__main__":
340     import argparse
341     parser = argparse.ArgumentParser()
342     parser.add_argument("--data_path", type=str, default="Homework 1/Data/
343                           BRCA_data.csv")
344     parser.add_argument("--outdir", type=str, default="Homework 1/Latex/
345                           Results/Problem_2")
346     parser.add_argument("--seed", type=int, default=0)

```

```
345     args = parser.parse_args()  
346  
347     run_problem2_pipeline(  
348         data_path=args.data_path,  
349         outdir=args.outdir,  
350         seed=args.seed  
351     )
```